
utils

HuangFuSL

Jan 27, 2026

CONTENTS

1	utils documentation	1
1.1	Usage	1
2	Polars-based Dataset Utilities	3
3	PyTorch Utilities	5
3.1	Automatic Device Selection	5
3.2	Functional Operators	6
3.3	Additional Metrics	9
3.4	Additional Neural Network Layers	11
3.5	Additional Base Operators	24
3.6	Utilities for Variable Length Sequences	25
3.7	Utilities for Reinforcement Learning	30
3.8	Utilities for Uplift Modeling	43
4	Sampler Utilities	47
5	Logging Utilities	49
6	Printing Utilities	51
7	Argument Parsing Utilities	53
	Python Module Index	57
	Index	59

CHAPTER
ONE

UTILS DOCUMENTATION

1.1 Usage

First mount this repo as a submodule under the source tree of your project. For example if your project is in `src/` directory.

```
git submodule add https://github.com/HuangFuSL/utils src/utils
```

Then you can import the modules in your Python code:

```
from utils import cprint
```

Use the following script to update the submodule:

```
git submodule update --remote --merge
```

**CHAPTER
TWO**

POLARS-BASED DATASET UTILITIES

PYTORCH UTILITIES

3.1 Automatic Device Selection

`utils.ctorch.device` - Utilities for managing and monitoring GPU devices in PyTorch.

`class utils.ctorch.device.GpuStat(idx, name, avg_util, avg_free_gb, total_gb)`

Bases: `object`

Dataclass that represents the status of a GPU device.

Parameters

- `idx (int)` – Index of the GPU.
- `name (str)` – Name of the GPU.
- `avg_util (float)` – Average utilization percentage over the sampling period.
- `avg_free_gb (float)` – Average free memory in GB over the sampling period.
- `total_gb (float)` – Total memory in GB of the GPU.

`avg_free_gb: float`

`avg_util: float`

`idx: int`

`name: str`

`total_gb: float`

`utils.ctorch.device.get_best_device(window_sec=3.0, interval_sec=0.5)`

Automatically selects the best available device based on GPU utilization and memory.

Parameters

- `window_sec (float)` – Total time in seconds to collect GPU statistics.
- `interval_sec (float)` – Time in seconds between each sample.

Returns

The best device identifier, either ‘mps’, ‘cuda:<idx>’, or ‘cpu’.

Return type

`str`

3.2 Functional Operators

`utils.ctorch.functional` - Functional utilities for PyTorch tensors.

`utils.ctorch.functional.gradient_reversal(x, alpha=1.0)`

Apply a gradient reversal layer to the input tensor. The forward pass is the identity function, but during back-propagation, the gradient is multiplied by -alpha.

Parameters

- `x (torch.Tensor)` – Input tensor.
- `alpha (float)` – Scaling factor for the gradient reversal. Default is 1.0.

Returns

The input tensor with the gradient reversed during backpropagation.

Return type

`torch.Tensor`

`utils.ctorch.functional.log_norm_pdf(x, mean, Sigma=None, logSigma=None, batch_first=None)`

Calculate the log probability density function of a normal distribution.

$$\log p(x) = -\frac{1}{2} \left(D \log(2\pi) + \log |\Sigma| + (x - \mu)^T \Sigma^{-1} (x - \mu) \right)$$

Parameters

- `x (torch.Tensor)` – Input tensor, shape (N, D), where N is the number of samples and D is the number of dimensions.
- `mean (torch.Tensor)` – Mean of the normal distribution, shape (D,), or (N, D)
- `Sigma (torch.Tensor / None)` – Covariance matrix of the normal distribution, shape (N, D, D), (N, D), (N,), (D, D), (D,), or a scalar.
- `logSigma (torch.Tensor / None)` – Logarithm of the covariance matrix, same shape as Sigma.
- `batch_first (bool / None)` – If True, indicates that the first dimension of Sigma is the batch size (N).

Returns

Tensor containing the log PDF values.

Return type

`torch.Tensor`

`utils.ctorch.functional.logit_product(x, y)`

Calculate the logit of the product of two probabilities given their logits.

$$\sigma^{-1}(\sigma(x) \cdot \sigma(y)) = x + y - \log(1 + e^x + e^y)$$

Parameters

- `x (torch.Tensor)` – Logit of the first probability.
- `y (torch.Tensor)` – Logit of the second probability.

Returns

Logit of the product of the two probabilities.

Return type

`torch.Tensor`

`utils.ctorch.functional.mmd_distance(x, y, *(Keyword-only parameters separator (PEP 3102)),
sigma=None, gamma=None, reduce=False)`

Compute the Maximum Mean Discrepancy (MMD) distance between two sets of tensors.

The MMD distance is given by:

$$\text{MMD}(x, y) = K(x, x) - 2K(x, y) + K(y, y)$$

Parameters

- **x** (`torch.Tensor`) – First tensor, shape (N, D), where N is the number of samples and D is the number of features.
- **y** (`torch.Tensor`) – Second tensor, shape (M, D), where M is the number of samples and D is the number of features.
- **sigma** (`torch.Tensor / int / float / None`) – Bandwidth parameter for the RBF kernel, scalar, or shape (K,), where K is the number of kernels.
- **gamma** (`torch.Tensor / int / float / None`) – $1 / (2 * \text{sigma}^2)$ parameter for the RBF kernel, scalar, or shape (K,), where K is the number of kernels.
- **reduce** (`bool / torch.Tensor`) – Whether to reduce the output. * If True, returns the mean MMD distance under different bandwidths. * If False, returns the MMD distance for each bandwidth. * If a tensor, it should have shape (K,) and will be used as mean weight.

Returns

Tensor containing the MMD distance values, shape (K,) if reduce is False, or a scalar otherwise,

Return type

`torch.Tensor`

`utils.ctorch.functional.norm_pdf(x, mean, Sigma=None, logSigma=None)`

Calculate the log probability density function of a normal distribution.

Parameters

- **x** (`torch.Tensor`) – Input tensor, shape (N, D), where N is the number of samples and D is the number of dimensions.
- **mean** (`torch.Tensor`) – Mean of the normal distribution, shape (D,), or (N, D)
- **Sigma** (`torch.Tensor / None`) – Covariance matrix of the normal distribution, shape (N, D, D), (N, D), (N,), (D, D), (D,), or a scalar.
- **logSigma** (`torch.Tensor / None`) – Logarithm of the covariance matrix, same shape as Sigma.
- **batch_first** (`bool`) – If True, indicates that the first dimension of Sigma is the batch size (N).

Returns

Tensor containing the PDF values.

Return type

`torch.Tensor`

`utils.ctorch.functional.rbf_kernel(x, y=None, *, sigma=None, gamma=None, reduce=False)`

Compute the Radial Basis Function (RBF) kernel between two sets of tensors.

The RBF kernel is given by:

$$K(x, y) = \exp(-\gamma ||x - y||^2)$$

or equivalently,

$$K(x, y) = \exp\left(-\frac{\|x - y\|^2}{2\sigma^2}\right)$$

Parameters

- **x** (`torch.Tensor`) – First tensor, shape (N, D), where N is the number of samples and D is the number of features.
- **y** (`torch.Tensor / None`) – Second tensor, shape (M, D), where M is the number of samples and D is the number of features.
- **sigma** (`torch.Tensor / int / float / None`) – Bandwidth parameter for the RBF kernel, scalar, or shape (K,), where K is the number of kernels.
- **gamma** (`torch.Tensor / int / float / None`) – $1 / (2 * \text{sigma}^2)$ parameter for the RBF kernel, scalar, or shape (K,), where K is the number of kernels.
- **reduce** (`torch.Tensor / bool`) – Whether to reduce the output.
 - If True, returns the mean of RBF kernel values under different bandwidths.
 - If False, returns the RBF kernel values for each bandwidth.
 - If a tensor, it should have shape (K,) and will be used as mean weight.

Returns

Tensor containing the RBF kernel values, shape (N, M) or (K, N, M) if multiple kernels are used.

Return type

`torch.Tensor`

```
utils.ctorch.functional.wasserstein_distance(x, y, p=2.0, eps=1e-06, wasser_iters=20,  
                                             wasser_eps=0.001)
```

Compute the Wasserstein distance between two sets of tensors using the Sinkhorn algorithm.

The Wasserstein distance is given by:

$$W_p(x, y) = \left(\inf_{\gamma \in \Gamma(x, y)} \int \|x - y\|^p d\gamma(x, y) \right)^{1/p}$$

Parameters

- **x** (`torch.Tensor`) – First tensor, shape (N, D), where N is the number of samples and D is the number of features.
- **y** (`torch.Tensor`) – Second tensor, shape (M, D), where M is the number of samples and D is the number of features.
- **p** (`float`) – Order of the norm to use for the distance calculation.
- **eps** (`float`) – Small value to avoid division by zero.
- **wasser_iters** (`int`) – Number of iterations for the Sinkhorn algorithm.
- **wasser_eps** (`float`) – Epsilon value for the Sinkhorn algorithm.

Returns

Tensor containing the Wasserstein distance value.

Return type

`torch.Tensor`

3.3 Additional Metrics

`utils.ctorch.metrics` - Utility functions for computing metrics in PyTorch

`class utils.ctorch.metrics.BatchedAUC(nbins=1000, device='cpu', logit=False)`

Bases: `BatchedMetric`

Batched AUC metric for binary classification tasks. This class computes the AUC in a batch-wise manner using a generator.

```
auc_metric = BatchedAUC(nbins=1000, device='cpu', logit=False)
for batch in data_loader:
    y_true, y_score = batch
    auc_metric(y_true, y_score)
auc_value = auc_metric.finalize()
```

`accumulator(nbins=1000, device='cpu', logit=False)`

Compute the AUC in a batch-wise manner using a generator.

Parameters

- `nbins` (`int`) – Number of bins to use for the histogram.
- `device` (`str` / `torch.device`) – Device to use for computation.
- `logit` (`bool`) – If True, apply sigmoid to the scores before computing AUC.

`class utils.ctorch.metrics.BatchedHitRate(k=None)`

Bases: `BatchedMetric`

`accumulator(k=None)`

Compute the hit rate in a batch-wise manner using a generator.

Parameters

- `k` (`int`) – The number of top predictions to consider.

`class utils.ctorch.metrics.BatchedMetric(**kwargs)`

Bases: ABC

`abstractmethod accumulator()`

Create a generator that accumulates metric values from batches of (`y_true`, `y_score`).

Return type

`Generator[None, Optional[Tuple[Tensor, Tensor]], float]`

Returns

Generator that yields None to receive batches and returns the computed metric value.

`finalize()`

Finalize the metric computation and return the accumulated value.

Returns

The computed metric value.

Return type

`float`

`reset()`

Reset the metric accumulator to start a new computation.

```
class utils.ctorch.metrics.BatchedNDCG(k=None)
```

Bases: *BatchedMetric*

accumulator(*k*=None)

Compute the NDCG in a batch-wise manner using a generator.

Parameters

k (*int*) – The number of top predictions to consider.

```
class utils.ctorch.metrics.MetricFormatter(name, starting_epoch=0, larger_better=True, eps=0.0005)
```

Bases: *object*

Formats the output of metrics during training.

format()

Format the metric for display.

Returns

Formatted string representation of the metric.

Return type

str

update(*value*)

Update the metric with a new value.

Parameters

value (*float*) – The new value to update the metric with.

```
utils.ctorch.metrics.auc_score(y_true, y_score)
```

Compute the Area Under the Curve (AUC) for binary classification.

Parameters

- **y_true** (*torch.Tensor*) – Ground truth binary labels (0 or 1).
- **y_score** (*torch.Tensor*) – Predicted scores or probabilities.

Returns

The computed AUC value.

Return type

float

```
utils.ctorch.metrics.auuc_score(y, t, tau, qini=True, normalize=True)
```

Compute the Area Under the Uplift Curve (AUUC).

Parameters

- **y** (*torch.Tensor*) – Binary outcome (0 or 1).
- **t** (*torch.Tensor*) – Binary treatment assignment (0 or 1).
- **tau** (*torch.Tensor*) – Uplift scores.
- **qini** (*bool*) – If True, compute Qini instead of AUUC (subtracts the random baseline).

Returns

The computed AUUC value.

Return type

float

```
utils.ctorch.metrics.hit_rate(y_true, y_score, k=None)
```

Compute the hit rate for retrieval tasks.

Parameters

- **y_true** (`torch.Tensor`) – Shape (N, C), ground truth binary labels (0 or 1).
- **y_score** (`torch.Tensor`) – Shape (N, C), predicted scores or probabilities.
- **k** (`int`) – The number of top predictions to consider.

Returns

The computed hit rate.

Return type

`float`

```
utils.ctorch.metrics.ndcg_score(y_true, y_score, k=None)
```

Compute the normalized discounted cumulative gain (NDCG) for retrieval tasks.

Parameters

- **y_true** (`torch.Tensor`) – Shape (N, C), ground truth binary labels (0 or 1).
- **y_score** (`torch.Tensor`) – Shape (N, C), predicted scores or probabilities.
- **k** (`int`) – The number of top predictions to consider.

Returns

The computed NDCG value.

Return type

`float`

3.4 Additional Neural Network Layers

nn.py - Utilities Modules for PyTorch tensors

Originally in ctorch.py

```
class utils.ctorch.nn.Activation(name, *args, **kwargs)
```

Bases: `Module`

Arbitrary activation function module.

Parameters

- **name** (`str`) – The name of the activation function.
- ***args** – Positional arguments for the activation function.
- ****kwargs** – Keyword arguments for the activation function.

forward(*x*)

Forward pass for the activation function.

Return type

`Tensor`

```
class utils.ctorch.nn.CholeskyTrilLinear(in_features, out_dim, bias=True, eps=0.0001, scale=None, non_neg_func='softplus')
```

Bases: `Module`

Implements a linear layer that returns a lower-triangular matrix that is positive definite.

Parameters

- **in_features** (*int*) – Number of input features.
- **out_dim** (*int*) – The output matrix dimension.
- **bias** (*bool*) – Whether to include a bias term.
- **eps** (*float*) – The small value added to the main diagonal of the matrix
- **scale** (*float*) – The maximum scale of the matrix elements.
- **non_neg_func** (*str* / *Callable*) – Element-wise non-negative activation function on the diagonal elements.
 - **softplus**: $f(x) = \log(1 + \exp(x))$
 - **elu**: $f(x) = ELU(x) + 1$
 - **sigmoid**: $f(x) = 1/(1 + e^{-x})$
 - **square**: $f(x) = x^2$
 - **exp**: $f(x) = e^x$

Shapes:

- Input shape: (*, in_features)
- Output shape: (*, out_dim, out_dim)

diag(*x*)

Return the main diagonal of the resulting matrix.

Parameters

x (*torch.Tensor*) – Input tensor of shape (*, in_features).

Return type

Tensor

Shapes:

- Input shape: (*, in_features)
- Output shape: (*, out_dim)

forward(*x*)

Forward pass for the positive definite linear layer.

Parameters

x (*torch.Tensor*) – Input tensor of shape (*, in_features).

Returns

Tril tensor of shape (*, out_dim, out_dim).

Return type

torch.Tensor

pd(*x*)

Return the positive definite LL^\top .

Parameters

x (*torch.Tensor*) – Input tensor of shape (*, in_features).

Return type

Tensor

Shapes:

- Input shape: $(*, \text{in_features})$
- Output shape: $(*, \text{out_dim}, \text{out_dim})$

```
class utils.ctorch.nn.DNN(*layer_dims, layer_type=<class 'torch.nn.modules.linear.Linear'>,  
                           flip_gradient=False, batchnorm=False, bias=True, dropout=None,  
                           activation='relu', residual=False)
```

Bases: *Module*

A Deep Neural Network (DNN) or Multi-Layer Perceptron (MLP) module.

Parameters

- **layer_dims** (*int) – The dimensions of each layer in the network, including input and output dimensions.
- **layer_type** (Type[*torch.nn.Module*]) – The type of layer to use (e.g., Linear).
- **flip_gradient** (bool) – Whether to apply a gradient reversal layer at the beginning.
- **batchnorm** (bool) – Whether to apply batch normalization after each linear layer.
- **bias** (bool) – Whether to include a bias term in the linear layers.
- **dropout** (float / None) – Dropout rate to apply after each layer. If None, no dropout is applied.
- **activation** (str / Activation / None) – Activation function to apply after each layer.
- **residual** (bool) – Whether to add a residual connection from input to output, requiring input and output dimensions to match.

Shapes:

- Input shape: $(*, \text{layer_dims}[0])$
- Output shape: $(*, \text{layer_dims}[-1])$

forward(*x*)

Forward pass for the DNN module.

Parameters

x (*torch.Tensor*) – Input tensor of shape $(*, \text{layer_dims}[0])$.

Returns

Output tensor of shape $(*, \text{layer_dims}[-1])$.

Return type

torch.Tensor

```
class utils.ctorch.nn.DeEmbedding(embedding)
```

Bases: *Module*

forward(*x*)

Forward pass for the de-embedding layer.

Parameters

x (*torch.Tensor*) – Tensor of shape $(*, \text{D})$, where D is the embedding dimension.

Returns

Tensor of shape $(*, \text{num_embeddings})$, where num_embeddings is the size of the embedding.

Return type

torch.Tensor

class utils.ctorch.nn.FactorizedNoisyLinear(*in_features*, *out_features*, *bias=True*, *init_sigma=0.5*)Bases: *Module*Implements a noisy linear layer according to <https://arxiv.org/abs/1706.10295>

The layer works the same way as a standard linear layer, but with added noise during training.

$$\begin{aligned} z_{\text{in}} &\sim \mathcal{N}(0, I_{d_{\text{in}}}) \\ z_{\text{out}} &\sim \mathcal{N}(0, I_{d_{\text{out}}}) \\ f(x) &= \text{sign}(x) \odot \sqrt{|x|} \\ w &= w_\mu + w_\sigma \odot (f(z_{\text{in}}) f(z_{\text{out}})^\top) \\ b &= b_\mu + b_\sigma \odot f(z_{\text{out}}) \end{aligned}$$

The parameters are initialized as:

$$\begin{aligned} w_\mu, b_\mu &\sim \mathcal{U}(-1/\sqrt{d_{\text{in}}}, 1/\sqrt{d_{\text{int}}}) \\ w_\sigma, b_\sigma &= \sigma_{\text{init}}/\sqrt{d_{\text{in}}} \end{aligned}$$

Parameters

- **in_features** (*int*) – Number of input features.
- **out_features** (*int*) – Number of output features.
- **bias** (*bool*) – Whether to include a bias term.
- **init_sigma** (*float*) – (*float*): The initial sigma coefficient σ_{init} , default is 0.5.

Shapes:

- Input shape: (*, *in_features*)
- Output shape: (*, *out_features*)

static *f*(*x*)**Return type**

Tensor

forward(*x*)

Forward pass for the noisy linear layer.

Parameters*x* (*torch.Tensor*) – Input tensor of shape (*, *in_features*).**Returns**Output tensor of shape (*, *out_features*).**Return type**

torch.Tensor

class utils.ctorch.nn.FeatureEmbedding(*num_features*, *embedding_size*, *padding_idx=None*, *max_norm=None*, *norm_type=2.0*, *scale_grad_by_freq=False*, *sparse=False*)Bases: *Module*

An embedding layer for encoding N multiple categorical features.

Parameters

- **num_features** (*List[int]*) – The number of unique values for each categorical feature.
- **embedding_size** (*List[int] / int*) – The size of the embedding for each feature. If a single integer is provided, it will be used for all features.
- **padding_idx** (*int / None*)
- **max_norm** (*float / None*)
- **norm_type** (*float*)
- **scale_grad_by_freq** (*bool*)
- **sparse** (*bool*)

Shapes:

- Input shape: (*, num_features)
- Output shape: (*, sum(embedding_size))

forward(*x*)

Forward pass for the embedding layer.

Parameters

x (*Tensor*) – Tensor of shape (*, num_features)

Returns

Tensor of shape (*, sum(embedding_size))

Return type

torch.Tensor

property total_embedding_size: int

Gets the total embedding size, which is the sum of all individual embedding sizes.

Returns

Total embedding size.

Return type

int

```
class utils.ctorch.nn.GaussianLinear(in_features, out_features, bias=True, mu_scale=None,
                                      Sigma_scale=None, eps=0.0001, non_neg_func='softplus',
                                      gaussian_type=<class
                                      'torch.distributions.multivariate_normal.MultivariateNormal'>)
```

Bases: *Module*

Implements a linear layer that returns a multivariate Gaussian distribution

Parameters

- **in_features** (*int*) – Number of input features.
- **out_features** (*int*) – Number of output features.
- **bias** (*bool*) – Whether to include a bias term.
- **mu_scale** (*float / None*) – If not None, the mean value is scaled by a tanh function with the given scale.
- **cov_scale** (*float*) – The boundary value for covariance matrix elements.
- **eps** (*float*) – Minimum value added to the diagonal of the covariance matrix.

- **non_neg_func** (*str* / *Callable*) – Non-negative mapping of the diagonal elements of the covariance matrix.
- **gaussian_type** ([MultivariateNormalClass](#)) – The constructor of target distribution.

Shapes:

- Input shape: (*, in_features)
- Output shape: (*, out_features)

cov(*x*)

Return the covariance matrix of the target distribution.

Parameters

x (*torch.Tensor*) – Input tensor of shape (*, in_features).

Return type

Tensor

Shapes:

- Input shape: (*, in_features)
- Output shape: (*, out_features, out_features)

forward(*x*)

Forward pass for the positive definite linear layer.

Parameters

x (*torch.Tensor*) – Input tensor of shape (*, in_features).

Returns

The target distribution.

Return type

torch.distributions.Distribution

mean(*x*)

Return the mean value of the target distribution.

Parameters

x (*torch.Tensor*) – Input tensor of shape (*, in_features).

Return type

Tensor

Shapes:

- Input shape: (*, in_features)
- Output shape: (*, out_features)

class *utils.ctorch.nn.GradientReversalLayer*(*alpha=1.0*)

Bases: *Module*

A layer that reverses the gradient during backpropagation.

Parameters

alpha (*float*) – The scaling factor for the gradient reversal. Default is 1.0.

forward(*x*)

Forward pass for the gradient reversal layer.

Parameters

x (`torch.Tensor`) – Input tensor.

Return type

Tensor

```
class utils.ctorch.nn.IndependentNoisyLinear(in_features, out_features, bias=True, init_sigma=0.017)
```

Bases: `Module`

Implements a noisy linear layer according to <https://arxiv.org/abs/1706.10295>

The layer works the same way as a standard linear layer, but with added noise during training.

$$\begin{aligned} w &= w_\mu + w_\sigma \odot \varepsilon \\ b &= b_\mu + b_\sigma \odot \varepsilon \end{aligned}$$

The parameters are initialized as:

$$\begin{aligned} w_\mu, b_\mu &\sim \mathcal{U}(-1/\sqrt{d_{\text{in}}}, 1/\sqrt{d_{\text{in}}}) \\ w_\sigma, b_\sigma &= \sigma_{\text{init}} \end{aligned}$$

Parameters

- **in_features** (`int`) – Number of input features.
- **out_features** (`int`) – Number of output features.
- **bias** (`bool`) – Whether to include a bias term.
- **init_sigma** (`float`) – (`float`): The initial sigma coefficient σ_{init} , default is 0.017.

Shapes:

- Input shape: $(*, \text{in_features})$
- Output shape: $(*, \text{out_features})$

forward(x)

Forward pass for the noisy linear layer.

Parameters

x (`torch.Tensor`) – Input tensor of shape $(*, \text{in_features})$.

Returns

Output tensor of shape $(*, \text{out_features})$.

Return type

`torch.Tensor`

```
class utils.ctorch.nn.LogStackedTruncatedNormal(lb=-inf, ub=inf, eps=1e-06,
                                                sigma_activation='softplus')
```

Bases: `StackedTruncatedNormal`

Implements the truncated log-normal distribution. The truncated probability density function is **stacked to the span boundaries**. The lb and ub are defined in the log-space.

Parameters

- **lb** (`float`) – The lower bound of the truncation in log-space, default is -inf.
- **ub** (`float`) – The upper bound of the truncation in log-space, default is inf.
- **eps** (`float`) – A small value to avoid numerical issues.

- **sigma_activation** (*str*) – The activation function to ensure positivity of sigma. Should be either ‘softplus’ or ‘exp’.

forward(*x*)

Define the computation performed at every call.

Should be overridden by all subclasses.

Note

Although the recipe for forward pass needs to be defined within this function, one should call the [Module](#) instance afterwards instead of this since the former takes care of running the registered hooks while the latter silently ignores them.

Return type

Tensor

loss(*x, target*)**Return type**

Tensor

class `utils.ctorch.nn.Module(*args, **kwargs)`

Bases: `Module`

A base class for all modules in `ctorch`. Supports device tracking and parameter counting.

debug()

Enable debug mode (only for custom `ctorch.nn.Modules`).

property device

Get the device of the module.

Returns

The device on which the module’s parameters are located.

Return type

`torch.device`

property num_parameters

Get the number of parameters in the module.

Returns

The total number of parameters in the module.

Return type

int

class `utils.ctorch.nn.MonotonicLinear(in_features, out_features, bias=True, non_neg_func='softplus')`

Bases: `Module`

Implements a monotonic linear layer. The monotonicity is enforced by applying a non-negative activation function to the weights.

Parameters

- **in_features** (*int*) – Number of input features.
- **out_features** (*int*) – Number of output features.

- **bias** (*bool*) – Whether to include a bias term.
- **non_neg_func** (*str* / *Callable*) – Element-wise non-negative activation function to use. Should be one of:
 - **relu**: $f(x) = \max(0, x)$
 - **softplus**: $f(x) = \log(1 + \exp(x))$
 - **sigmoid**: $f(x) = \frac{1}{1+\exp(-x)}$
 - **elu**: $f(x) = ELU(x) + 1$
 - **abs**: $f(x) = |x|$
 - **square**: $f(x) = x^2$
 - **exp**: $f(x) = e^x$

To keep the monotonicity, non-monotonic activations including `softmax`, `GELU`, `SiLU` and `Mish` should not be used. Normalization techniques such as layer normalization or batch normalization should also be avoided.

Shapes:

- Input shape: $(*, \text{in_features})$
- Output shape: $(*, \text{out_features})$

`forward`(*x*)

Forward pass for the monotonic linear layer.

Parameters

x (*torch.Tensor*) – Input tensor of shape $(*, \text{in_features})$.

Returns

Output tensor of shape $(*, \text{out_features})$.

Return type

torch.Tensor

`class utils.ctorch.nn.MultivariateNormalClass(*args, **kwargs)`

Bases: *Protocol*

Internal protocol for `GaussianLinear` module. Any valid `gaussian_type` parameter should follow the following protocol.

```
def __call__(
    self, loc: torch.Tensor, covariance_matrix: torch.Tensor | None = None,
    precision_matrix: torch.Tensor | None = None,
    scale_tril: torch.Tensor | None = None,
    validate_args: Any = None
) -> torch.distributions.Distribution:
    ...
```

`class utils.ctorch.nn.NegativeBinomial(alpha_param=False, activation='softplus')`

Bases: *Module*

Implements the negative log-likelihood loss for the Negative Binomial distribution.

Let g_μ and g_α be the outputs of the network. The mean μ and dispersion α are obtained via `activation` function to ensure positivity.

The log-likelihood loss is computed as:

$$\begin{aligned}\mathcal{L}(x; \mu, \alpha) = & \log \Gamma(x + \alpha) - \log \Gamma(\alpha) - \log \Gamma(x + 1) \\ & + \alpha(\log \alpha - \log(\mu + \alpha)) + x(\log \mu - \log(\mu + \alpha))\end{aligned}$$

Parameters

- **alpha_param** (*bool*) – If True, use a learnable parameter for alpha. Default is False.
- **activation** (*str*) – The activation function to ensure positivity of mu and alpha. Should be either ‘softplus’ or ‘exp’. Default is ‘softplus’.

Shapes:

- Input shape: (*, 2) if alpha_param is False, else (*)
- Output shape: (*)

forward(*x*)

Define the computation performed at every call.

Should be overridden by all subclasses.

Note

Although the recipe for forward pass needs to be defined within this function, one should call the *Module* instance afterwards instead of this since the former takes care of running the registered hooks while the latter silently ignores them.

Return type

Tensor

loss(*x*, *target*)

Return type

Tensor

class `utils.ctorch.nn.RotaryTemporalEmbedding(embedding_dim, denom=10000.0)`

Bases: *Module*

Implements rotary positional embedding proposed in “RoFormer: Enhanced Transformer with Rotary Position Embedding” (<https://arxiv.org/abs/2104.09864>).

$$\begin{aligned}\mathbf{R} &= \text{diag}(\mathbf{R}_1, \dots, \mathbf{R}_{\lfloor n/2 \rfloor}) \\ \mathbf{R}_i &= \begin{bmatrix} \cos(t\theta_i) & -\sin(t\theta_i) \\ \sin(t\theta_i) & \cos(t\theta_i) \end{bmatrix} \\ \theta_i &= \frac{1}{10000^{2(i-1)/d_{model}}}\end{aligned}$$

Parameters

- **embedding_dim** (*int*) – The dimension of the embedding space. Must be even.
- **denom** (*float*) – The denominator (10000.0) for the positional encoding.

Shapes:

- Input shape: x (*, embedding_dim), t (*)

- Output shape: (*, embedding_dim)

forward(*t, x*)

Forward pass for the rotary temporal embedding.

Parameters

- ***t*** (*torch.Tensor*) – Time record tensor of shape (*).
- ***x*** (*torch.Tensor*) – Input tensor of shape (*, embedding_dim).

Returns

Embedding of shape (*, embedding_dim).

Return type

torch.Tensor

class utils.ctorch.nn.SinusoidalTemporalEmbedding(*embedding_dim, denom=10000.0*)

Bases: *Module*

Implements sinusoidal positional embedding proposed in “Attention is All You Need”.

$$PE_{(batch, pos, i)} = \begin{cases} \sin\left(\frac{pos}{10000^{2k/d_{model}}}\right) & \text{if } i = 2k \\ \cos\left(\frac{pos}{10000^{2k/d_{model}}}\right) & \text{if } i = 2k + 1 \end{cases}$$

Parameters

- ***embedding_dim*** (*int*) – The dimension of the embedding space.
- ***denom*** (*float*) – The denominator (10000.0) for the positional encoding.

Shapes:

- Input shape: (*, embedding_dim)
- Output shape: (*, embedding_dim)

forward(*t*)

Forward pass for the circular temporal embedding.

Parameters

t (*torch.Tensor*) – Time record tensor of shape (*).

Returns

Embedding of shape (*, embedding_dim).

Return type

torch.Tensor

class utils.ctorch.nn.StackedTruncatedNormal(*lb=-inf, ub=inf, eps=1e-06, sigma_activation='softplus'*)

Bases: *Module*

Implements the truncated normal distribution. The truncated probability density function is **stacked to the span boundaries**.

Parameters

- ***lb*** (*float*) – The lower bound of the truncation, default is -inf.
- ***ub*** (*float*) – The upper bound of the truncation, default is inf.
- ***eps*** (*float*) – A small value to avoid numerical issues.

- **sigma_activation** (str) – The activation function to ensure positivity of sigma. Should be either ‘softplus’ or ‘exp’.

forward(*x*)

Define the computation performed at every call.

Should be overridden by all subclasses.

Note

Although the recipe for forward pass needs to be defined within this function, one should call the [Module](#) instance afterwards instead of this since the former takes care of running the registered hooks while the latter silently ignores them.

Return type

Tensor

loss(*x*, *target*)**Return type**

Tensor

```
class utils.ctorch.nn.TransformerDecoderLayer(d_model, nhead, dim_feedforward=2048, dropout=0.1,
                                              activation=<function relu>, layer_norm_eps=1e-05,
                                              batch_first=False, norm_first=False, bias=True,
                                              device=None, dtype=None)
```

Bases: TransformerDecoderLayer

```
get_cross_attention_map(tgt, memory, tgt_mask=None, memory_mask=None,
                        tgt_key_padding_mask=None, memory_key_padding_mask=None,
                        tgt_is_causal=False, memory_is_causal=False)
```

Get the cross-attention map from the decoder layer.

Parameters

- **tgt** (`torch.Tensor`)
- **memory** (`torch.Tensor`)
- **tgt_mask** (`torch.Tensor` / `None`)
- **memory_mask** (`torch.Tensor` / `None`)
- **tgt_key_padding_mask** (`torch.Tensor` / `None`)
- **memory_key_padding_mask** (`torch.Tensor` / `None`)
- **tgt_is_causal** (`bool`)
- **memory_is_causal** (`bool`)

Returns

The attention weights of shape (batch_size, num_heads, tgt_seq_len, memory_seq_len).

Return type

`torch.Tensor`

```
get_self_attention_map(tgt, memory, tgt_mask=None, memory_mask=None,
                      tgt_key_padding_mask=None, memory_key_padding_mask=None,
                      tgt_is_causal=False, memory_is_causal=False)
```

Get the self-attention map from the decoder layer.

Parameters

- **tgt** (`torch.Tensor`)
- **tgt_mask** (`torch.Tensor / None`)
- **memory_mask** (`torch.Tensor / None`)
- **tgt_key_padding_mask** (`torch.Tensor / None`)
- **memory_key_padding_mask** (`torch.Tensor / None`)
- **tgt_is_causal** (`bool`)
- **memory_is_causal** (`bool`)

Returns

The attention weights of shape (batch_size, num_heads, seq_len, seq_len).

Return type

`torch.Tensor`

```
class utils.ctorch.nn.TransformerEncoderLayer(d_model, nhead, dim_feedforward=2048, dropout=0.1,
                                              activation=<function relu>, layer_norm_eps=1e-05,
                                              batch_first=False, norm_first=False, bias=True,
                                              device=None, dtype=None)
```

Bases: `TransformerEncoderLayer`

A Transformer encoder layer with additional functionality to get attention maps.

get_attention_map(`src, src_mask=None, src_key_padding_mask=None, is_causal=False`)

Get the attention map from the encoder layer.

Parameters

- **src** (`torch.Tensor`)
- **src_mask** (`torch.Tensor / None`)
- **src_key_padding_mask** (`torch.Tensor / None`)
- **is_causal** (`bool`)

Returns

The attention weights of shape (batch_size, num_heads, seq_len, seq_len).

Return type

`torch.Tensor`

```
class utils.ctorch.nn.ZeroInflatedLogNormal(zero=0.0, eps=1e-08)
```

Bases: `Module`

Implements a zero-inflated log-normal distribution layer. A zero-inflated log-normal distribution is a mixture of a point mass at zero and a log-normal distribution.

$$\begin{aligned} P(X = 0) &= \sigma(-\text{logit}) \\ P(X > 0) &= 1 - \sigma(-\text{logit}) \\ \log X | X > 0 &\sim \mathcal{N}(\mu, \sigma^2) \end{aligned}$$

Parameters

- **zero** (`float`) – The minimum value for the output.

- **eps** (*float*) – A small value to avoid numerical issues.

Shapes:

- Input shape: $(*, 2)$ or $(*, 3)$, where the last dimension represents (logit, mu) or (logit, mu, log_sigma).
- Output shape: $(*)$

check_input_shape(*x*)

Check if the input tensor has sigma dimension.

Parameters

x (*torch.Tensor*) – Input tensor of shape $(*, 2)$ or $(*, 3)$.

Returns

The (logit, mu, log_sigma) tensors.

Return type

Tuple[torch.Tensor, ...]

forward(*x*)

Forward pass for the zero-inflated log-normal distribution.

Parameters

x (*torch.Tensor*) – Input tensor of shape $(*, 2)$ or $(*, 3)$.

Returns

Output tensor of shape $(*)$.

Return type

torch.Tensor

loss(*x, target*)

Compute the negative log-likelihood loss for the zero-inflated log-normal distribution.

Parameters

- **x** (*torch.Tensor*) – Input tensor of shape $(*, 2)$ or $(*, 3)$.

- **target** (*torch.Tensor*) – Target tensor of shape $(*)$.

Returns

Loss tensor of shape $(*)$.

Return type

torch.Tensor

3.5 Additional Base Operators

utils.ctorch.ops - Utilities Operators for PyTorch tensors

class utils.ctorch.ops.GradientReversalOp(*args, **kwargs)

Bases: Function

Gradient reversal operation for adversarial training.

static backward(*ctx, grad_output*)

Backward pass for the gradient reversal operation.

Parameters

grad_output (*torch.Tensor*) – Gradient from the next layer.

Returns

The gradient with the reversal applied.

Return type

`torch.Tensor`

static forward(ctx, x, alpha)

Forward pass for the gradient reversal operation.

Parameters

- **x** (`torch.Tensor`) – Input tensor.
- **alpha** (`float`) – Scaling factor for the gradient reversal.

Returns

The input tensor with the gradient reversal applied.

Return type

`torch.Tensor`

3.6 Utilities for Variable Length Sequences

`padding.py` - Utilities for handling PackedSequences

Originally in `torch.py` Author: HuangFuSL Date: 2025-06-26

utils.ctorch.padding.append_right(input_tensor, append_value=0.0, num_steps=1, batch_first=True)

Add `num_steps` steps at the end of the time axis, to a packed sequence or tensor.

Parameters

- **input_tensor** (`PackedOrTensor`) – The input tensor or packed sequence to pad.
- **padding_value** (`float, optional`) – The value to use for padding. Defaults to 0.0.
- **num_steps** (`int, optional`) – The number of steps to add. Defaults to 1.
- **batch_first** (`bool, optional`) – If True, the input tensor is assumed to be in (batch, seq, ...) format, otherwise (seq, batch, ...), only effective when accepting tensors. Defaults to True.

Returns

The padded tensor or packed sequence.

Return type

`PackedOrTensor`

utils.ctorch.padding.get_key_padding_mask_left(lengths)

Create a key padding mask for sequences based on their lengths, with sequences left-aligned.

Parameters

- **lengths** (`torch.Tensor`) – A 1D tensor containing the lengths of each sequence.

Returns

A boolean mask tensor indicating the padding positions.

Return type

`torch.Tensor`

utils.ctorch.padding.get_key_padding_mask_right(lengths)

Create a key padding mask for sequences based on their lengths, with sequences right-aligned.

Parameters

lengths (`torch.Tensor`) – A 1D tensor containing the lengths of each sequence.

Returns

A boolean mask tensor indicating the padding positions.

Return type

`torch.Tensor`

`utils.ctorch.padding.get_model_memory_size(model)`

Get the total memory size of a model in bytes.

Parameters

model (`torch.nn.Module`) – The input model.

Returns

The total memory size of the model in bytes.

Return type

`int`

`utils.ctorch.padding.get_tensor_memory_size(tensor)`

Get the memory size of a tensor in bytes.

Parameters

tensor (`torch.Tensor`) – The input tensor.

Returns

The memory size of the tensor in bytes.

Return type

`int`

`utils.ctorch.padding.masked_select(values_input, mask_input)`

Perform masked selection of variable length on a PackedSequence or a regular tensor.

Parameters

- **values_input** (`PackedOrTensor`) – The input values to select from.
- **mask_input** (`PackedOrTensor`) – The mask indicating which values to select.

Returns

A PackedSequence containing only the selected values.

Return type

`PackedSequence`

Example

```
from torch.nn.utils.rnn import pad_packed_sequence
values_list = [
    [[1., 10.], [2., 20.], [3., 30.]],
    [[4., 40.], [5., 50.], [0., 0.]],
]
mask_list = [
    [True, False, True],
    [False, True, False],
]
```

(continues on next page)

(continued from previous page)

```

values = torch.tensor(values_list, dtype=torch.float32)
mask = torch.tensor(mask_list, dtype=torch.bool)

out = masked_select(values, mask)
out_padded, out_len = pad_packed_sequence(out, batch_first=True)

assert out_padded.tolist() == [
    [[1.0, 10.0], [3.0, 30.0]],
    [[5.0, 50.0], [0.0, 0.0]],
]
assert out_len.tolist() == [2, 1]

```

`utils.ctorch.padding.pack_padded_sequence_right(input, lengths, batch_first=False, enforce_sorted=False)`

Like `torch.nn.utils.rnn.pack_padded_sequence` but accepts right-aligned sequences.

Parameters

- **input** (`torch.Tensor`) – The input tensor, which should be right-aligned.
- **lengths** (`torch.Tensor`) – A 1D tensor containing the lengths of each sequence.
- **batch_first** (`bool`) – If True, the input is expected to be of shape (batch_size, seq_len, ...). Otherwise, the input is expected to be of shape (seq_len, batch_size, ...).
- **enforce_sorted** (`bool`) – If True, the input sequences must be sorted by length in descending order. If False, the input sequences can be in any order.

Returns

A packed sequence object containing the right-aligned sequences.

Return type

`torch.nn.utils.rnn.PackedSequence`

`utils.ctorch.padding.packed_binary_op(op, a, b)`

Apply a binary element-wise operation to a PackedSequence or a regular tensor.

Parameters

- **op** (`Callable`) – A binary operation to apply.
- **a** (`PackedSequence` / `torch.Tensor`) – The first input data, either a PackedSequence or a regular tensor.
- **b** (`PackedSequence` / `torch.Tensor`) – The second input data, either a PackedSequence or a regular tensor.

Returns

The output after applying the operation.

Return type

`PackedSequence` | `torch.Tensor`

`utils.ctorch.padding.packed_concat(packed_seq, dim=-1)`

Concatenate a list of PackedSequence objects along a specified dimension. Notice that the length of the packed sequences must be the same.

Parameters

- **packed_seq** (`List[PackedSequence]`) – List of PackedSequence objects to concatenate.

- **dim** (*int*) – Dimension along which to concatenate. Default is -1 (last dimension). The dimension must not be 0 (the packed time dimension). The sequence length dimension (dimension 1 of the padded tensor where batch_size is True) is omitted.

Returns

A new PackedSequence object containing the concatenated data.

Return type

PackedSequence

`utils.ctorch.padding.packed_forward(module, packed_input)`

Forward pass for a module with packed input.

Parameters

- **module** (`torch.nn.Module`) – The neural network to apply.
- **packed_input** (`PackedSequence` / `torch.Tensor`) – The packed input data.

Returns

The output after applying the module. If the input is a PackedSequence, the output will also be a PackedSequence, otherwise it will be a regular tensor.

Return type

PackedSequence | torch.Tensor

`utils.ctorch.padding.packed_unary_op(func, x)`

Apply an unary sample-wise function to a PackedSequence or a regular tensor.

Parameters

- **func** (`Callable`) – An sample-wise function to apply.
- **x** (`PackedSequence` / `torch.Tensor`) – The input data, either a PackedSequence or a regular tensor.

Returns

The output after applying the function. If the input is a PackedSequence, the output will also be a PackedSequence, otherwise it will be a regular tensor.

Return type

PackedSequence | torch.Tensor

`utils.ctorch.padding.pad_packed_sequence_right(sequence, batch_first=False, padding_value=0.0, total_length=None)`

Like `torch.nn.utils.rnn.pad_packed_sequence` but right-aligns the sequences.

Parameters

- **sequence** (`torch.nn.utils.rnn.PackedSequence`) – The packed sequence to pad.
- **batch_first** (`bool`) – If True, the output will be of shape (batch_size, seq_len, ...). If False, the output will be of shape (seq_len, batch_size, ...).
- **padding_value** (`float`) – The value to use for padding.
- **total_length** (`int` / `None`) – If specified, the output will be padded to this length. If None, the output will be padded to the maximum length of the sequences in the packed sequence

Returns

The padded tensor and the lengths of the original sequences.

Return type

Tuple[torch.Tensor, torch.Tensor]

`utils.ctorch.padding.prepend_left(input_tensor, append_value=0.0, num_steps=1, batch_first=True)`

Add `num_steps` steps at the beginning of the time axis, to a packed sequence or tensor.

Parameters

- **input_tensor** (*PackedOrTensor*) – The input tensor or packed sequence to pad.
- **padding_value** (*float, optional*) – The value to use for padding. Defaults to 0.0.
- **num_steps** (*int, optional*) – The number of steps to add. Defaults to 1.
- **batch_first** (*bool, optional*) – If True, the input tensor is assumed to be in (batch, seq, ...) format, otherwise (seq, batch, ...), only effective when accepting tensors. Defaults to True.

Returns

The padded tensor or packed sequence.

Return type

PackedOrTensor

`utils.ctorch.padding.truncate_left(input_tensor, num_steps=1, batch_first=True, strict=True)`

Remove `num_steps` steps at the beginning of the time axis, from a packed sequence or tensor.

Parameters

- **input_tensor** (*PackedOrTensor*) – The input tensor or packed sequence to truncate.
- **num_steps** (*int, optional*) – The number of steps to remove. Defaults to 1.
- **batch_first** (*bool, optional*) – If True, the input tensor is assumed to be in (batch, seq, ...) format, otherwise (seq, batch, ...), only effective when accepting tensors. Defaults to True.
- **strict** (*bool, optional*) – If True, raises an error if truncating action will yield an empty sequence.

Returns

The truncated tensor or packed sequence.

Return type

PackedOrTensor

`utils.ctorch.padding.truncate_right(input_tensor, num_steps=1, batch_first=True, strict=True)`

Remove `num_steps` steps at the end of the time axis, from a packed sequence or tensor.

Parameters

- **input_tensor** (*PackedOrTensor*) – The input tensor or packed sequence to truncate.
- **num_steps** (*int, optional*) – The number of steps to remove. Defaults to 1.
- **batch_first** (*bool, optional*) – If True, the input tensor is assumed to be in (batch, seq, ...) format, otherwise (seq, batch, ...), only effective when accepting tensors. Defaults to True.
- **strict** (*bool, optional*) – If True, raises an error if truncating action will yield an empty sequence.

Returns

The truncated tensor or packed sequence.

Return type

PackedOrTensor

`utils.ctorch.padding.unpad_sequence_right(input, lengths, batch_first=False)`

Like `torch.nn.utils.rnn.unpad_sequence` but accepts right-aligned sequences.

Parameters

- **input** (`torch.Tensor`) – The input tensor, which should be right-aligned.
- **lengths** (`torch.Tensor`) – A 1D tensor containing the lengths of each sequence.
- **batch_first** (`bool`) – If True, the input is expected to be of shape (batch_size, seq_len, ...). Otherwise, the input is expected to be of shape (seq_len, batch_size, ...).

Returns

A list of tensors, each representing a sequence with padding removed.

Return typeList[`torch.Tensor`]

3.7 Utilities for Reinforcement Learning

rl.py - Utilities Components for Reinforcement Learning

```
class utils.ctorch.rl.BaseDistributionalQNetwork(state_shape, num_actions, *, gamma=0.99,
                                                 tau=1)
```

Bases: `BaseQNetwork`, ABC

Abstract base class for Distributional DQN with discrete action space.

One should implement the `forward` method to return un-softmaxed logits or log-softmaxed logits for all actions or specified actions, in shape (... , num_actions, num_atoms) or (... , num_atoms).

Parameters

- **state_shape** (`Tuple[int, ...]`) – The shape dimensions of the input state.
- **num_actions** (`int`) – The number of actions the agent can take.
- **atoms** (`List[int | float] | torch.Tensor`) – The support of the value distribution.
- **gamma** (`float, optional`) – The discount factor for future rewards, defaults to 0.99.
- **tau** (`int, optional`) – The number of steps to look ahead for target updates, defaults to 1.

`Q(state, action)`

Compute the action-value function $Q(s, a)$ for a given state s and action a .

Parameters

- **state** (`torch.Tensor`) – The input state tensor.
- **action** (`torch.Tensor`) – The input action tensor.

Returns

The action-value function $Q(s, a)$ for a given state s and action a .

Return type`torch.Tensor`

Shapes:

- state: $(*, (\text{state_dims},))$
- action: $(*)$
- output: $(*)$

`Q_all(state)`

Compute the action-value function $Q(s, a)$ for all actions a given state s .

Parameters

state (`torch.Tensor`) – The input state tensor.

Returns

The action-value function $Q(s, a)$ for all actions a .

Return type

`torch.Tensor`

Shapes:

- state: $(*, (\text{state_dims},))$
- output: $(*, \text{num_actions})$

`Q_all_dist(state)`

Compute the action-value distribution $Q(s, a)$ for all actions a given state s .

Parameters

state (`torch.Tensor`) – The input state tensor.

Returns

The action-value distribution $Q(s, a)$ for all actions a .

Return type

`torch.Tensor`

Shapes:

- state: $(*, (\text{state_dims},))$
- output: $(*, \text{num_actions}, \text{num_atoms})$

`Q_dist(state, action)`

Compute the action-value distribution $Q(s, a)$ for a given state s and action a .

Parameters

- **state** (`torch.Tensor`) – The input state tensor.
- **action** (`torch.Tensor`) – The input action tensor.

Returns

The action-value distribution $Q(s, a)$ for a given state s and action a .

Return type

`torch.Tensor`

Shapes:

- state: $(*, (\text{state_dims},))$
- action: $(*)$

- output: (*, num_atoms)

V_dist(state)

Compute the state-value distribution $V(s)$ for a given state s .

Parameters

state (`torch.Tensor`) – The input state tensor.

Returns

The state-value distribution $V(s)$ for a given state s .

Return type

`torch.Tensor`

Shapes:

- state: (*, (state_dims,))
- output: (*, num_atoms)

class utils.ctorch.rl.BasePolicyNetwork(state_shape, *, gamma=0.99, tau=1)

Bases: `BaseRLModel`, `TargetNetworkMixin`

Abstract base class for policy-based reinforcement learning models.

Parameters

- **state_shape** (`Tuple[int, ...]`) – The shape dimensions of the input state.
- **gamma** (`float, optional`) – The discount factor for future rewards, defaults to 0.99.
- **tau** (`int, optional`) – The number of steps to look ahead for target updates, defaults to 1.

act(state)

Return the action sampled from the policy given the state.

Parameters

state (`torch.Tensor`) – The input state tensor.

Returns

A tuple containing the sampled action and its log probability.

Return type

`Tuple[torch.Tensor, torch.Tensor]`

cumulative_reward(r, lambda_=1)

Calculate the cumulative reward for a trajectory, or the GAE(lambda) estimate.

Parameters

- **r** (`torch.Tensor`) – The reward tensor of shape (L,), for GAE, r is the state-value TD error $r_t + \gamma V(s_{t+1}) - V(s_t)$.
- **lambda_** (`float, optional`) – The lambda parameter for GAE. Defaults to 1 for cumulative reward.

Returns

The cumulative reward or GAE estimate of shape (L,).

Return type

`torch.Tensor`

Shapes:

- r: (L,)
- output: (L,)

abstractmethod forward(state)

Forward pass to compute the action distribution given the state.

Parameters

state (`torch.Tensor`) – The input state tensor.

Returns

The action distribution given the state.

Return type

`torch.distributions.Distribution`

log_pi(state, action)

Return the log probability of the action given the state.

Return type

`Tensor`

normalize_trajectory(r, eps=1e-08)

Return type

`Tensor`

pi(state, action)

Return the action distribution probability given the state.

Return type

`Tensor`

policy_parameters()

Get the parameters of the policy network.

Returns

An iterator over the parameters of the policy network.

Return type

`Iterator[torch.nn.Parameter]`

setup_target()

Create a target network by copying the current network. This method can only be called once.

property value_model: `BaseValueNetwork`

Get the state value sub-network. If not implemented, raises `NotImplementedError`.

value_parameters()

Get the parameters of the value network.

Returns

An iterator over the parameters of the value network.

Return type

`Iterator[torch.nn.Parameter]`

```
class utils.ctorch.rl.BaseQNetwork(state_shape, num_actions, *, gamma=0.99, tau=1)
```

Bases: *BaseValueNetwork*, ABC

Abstract base class for DQN with discrete action space.

One should either implement the `forward` method to return Q-values for all actions or specified actions.

Parameters

- **state_shape** (`Tuple[int, ...]`) – The shape dimensions of the input state.
- **num_actions** (`int`) – The number of actions the agent can take.
- **gamma** (`float, optional`) – The discount factor for future rewards, defaults to 0.99.
- **tau** (`int, optional`) – The number of steps to look ahead for target updates, defaults to 1.

A_all(*state*)

Compute the advantage function $A(s, a)$ for a given state s and action a .

$$A(s, a) = Q(s, a) - V(s)$$

Parameters

state (`torch.Tensor`) – The input state tensor.

Returns

The advantage function $A(s, a)$ for a given state s and all actions.

Return type

`torch.Tensor`

Shapes:

- state: `(*, (state_dims,))`
- action: `(*)`
- output: `(*)`

Q_all(*state*)

Compute the action-value function $Q(s, a)$ for all actions a given state s .

Parameters

state (`torch.Tensor`) – The input state tensor.

Returns

The action-value function $Q(s, a)$ for all actions a .

Return type

`torch.Tensor`

Shapes:

- state: `(*, (state_dims,))`
- output: `(*, num_actions)`

V(*state*)

Compute the state-value function $V(s)$ for a given state s .

Parameters

state (`torch.Tensor`) – The input state tensor.

Returns

The state-value function $V(s)$ for a given state s .

Return type

`torch.Tensor`

Shapes:

- state: $(*, (\text{state_dims},))$
- output: $(*)$

`act(state, eps=0.0, sample_wise=True)`

Policy for selecting actions based on the current state and exploration rate. By default, the policy takes exploration actions with a probability of ε and exploitation actions with a probability of $1 - \varepsilon$.

Parameters

`state (torch.Tensor)` – The input state tensor.

Returns

The selected actions.

Return type

`torch.Tensor`

`exploit(state)`

The exploitation policy, selecting actions based on the current Q-values.

Parameters

`state (torch.Tensor)` – The input state tensor.

Returns

The selected actions.

Return type

`torch.Tensor`

`explore(state)`

The exploration policy, defaulting to random actions.

Parameters

`state (torch.Tensor)` – The input state tensor.

Returns

The selected actions.

Return type

`torch.Tensor`

`abstractmethod forward(state, action=None)`

Returns the Q-values for the given state and action.

Parameters

- `state (torch.Tensor)` – The input state tensor.
- `action (torch.Tensor / None)` – The input action tensor.

Returns

The Q-values for the given state and action. If action is `None`, returns the Q-value over all actions.

Return type`torch.Tensor`**td_step**(*trajectory*, *a_prime*=None)**class** `utils.ctorch.rl.BaseRLModel`(*state_shape*, *, *gamma*=0.99, *tau*=1)Bases: `Module`, ABC

Abstract base class for reinforcement learning models.

Parameters

- **state_shape** (`Tuple[int, ...]`) – The shape dimensions of the input state.
- **gamma** (`float, optional`) – The discount factor for future rewards, defaults to 0.99.
- **tau** (`int, optional`) – The number of steps to look ahead for target updates, defaults to 1.

abstractmethod **act**(*state*, ***kwargs*)Policy for selecting actions based on the current state and exploration rate. By default, the policy takes exploration actions with a probability of ε and exploitation actions with a probability of $1 - \varepsilon$.**Parameters****state** (`torch.Tensor`) – The input state tensor.**Returns**

A tuple containing the selected action and the log probability of the action.

Return type`Tuple[torch.Tensor, torch.Tensor]`**abstractmethod** **loss**()**Return type**`Tensor`**class** `utils.ctorch.rl.BaseValueNetwork`(*state_shape*, *, *gamma*=0.99, *tau*=1)Bases: `BaseRLModel`, `TargetNetworkMixin`

Abstract base class for Value Networks.

One should either implement the `forward` method, or `V` method.**Parameters**

- **state_shape** (`Tuple[int, ...]`) – The shape dimensions of the input state.
- **gamma** (`float, optional`) – The discount factor for future rewards, defaults to 0.99.
- **tau** (`int, optional`) – The number of steps to look ahead for target updates, defaults to 1.

A(*state*, *action*)Compute the advantage function $A(s, a)$ for a given state s and action a .

$$A(s, a) = Q(s, a) - V(s)$$

Parameters

- **state** (`torch.Tensor`) – The input state tensor.
- **action** (`torch.Tensor`) – The input action tensor.

Returns

The advantage function $A(s, a)$ for a given state s and action a .

Return type

`torch.Tensor`

Shapes:

- state: $(*, (\text{state_dims},))$
- action: $(*)$
- output: $(*)$

`Q(state, action)`

Compute the action-value function $Q(s, a)$ for a given state s and action a .

Parameters

- **state** (`torch.Tensor`) – The input state tensor.
- **action** (`torch.Tensor`) – The input action tensor.

Returns

The action-value function $Q(s, a)$ for a given state s and action a .

Return type

`torch.Tensor`

Shapes:

- state: $(*, (\text{state_dims},))$
- action: $(*, (\text{action_dims},))$ or $(*)$
- output: $(*)$

`V(state)`

Compute the state-value function $V(s)$ for a given state s .

Parameters

state (`torch.Tensor`) – The input state tensor.

Returns

The state-value function $V(s)$ for a given state s .

Return type

`torch.Tensor`

Shapes:

- state: $(*, (\text{state_dims},))$
- output: $(*)$

`act(state, **kwargs)`

Policy for selecting actions based on the current state and exploration rate. By default, the policy takes exploration actions with a probability of ε and exploitation actions with a probability of $1 - \varepsilon$.

Parameters

state (`torch.Tensor`) – The input state tensor.

Returns

A tuple containing the selected action and the log probability of the action.

Return type

Tuple[torch.Tensor, torch.Tensor]

action_td_step(trajectory, a_prime)

The loss function for training the value network using the TD loss. If a_prime is provided, the action-value function Q(s, a) is used. Otherwise, the state-value function V(s) is used.

$$\begin{aligned} \text{LHS} &= V(s) \\ \text{RHS} &= \begin{cases} r + V(s') & \text{if not terminal} \\ r & \text{if terminal} \end{cases} \end{aligned}$$

The LHS and RHS can be optimized using `MSELoss` or `SmoothL1Loss`.

Parameters

- **trajectory** (`Trajectory`) – A batch of trajectories.
- **a_prime** (`torch.Tensor / None`) – The next action tensor. If provided, the action-value function Q(s, a) is used. If not provided, the state-value function V(s) is used.

Returns

The LHS and RHS of the TD loss.

Return type

Tuple[torch.Tensor, torch.Tensor]

abstractmethod forward(state, action=None)

Forward pass of the value network. The implementation should either implement the state-value function V(s) or the action-value function Q(s, a).

Parameters

- **state** (`torch.Tensor`) – The input state tensor.
- **action** (`torch.Tensor / None`) – The input action tensor. If provided, the network should compute the action-value function Q(s, a). If not provided, the network should compute the state-value function V(s).

Returns

The value function V(s) for the given state.

Return type

`torch.Tensor`

Shapes:

- state: `(*, (state_dims,))`
- action: `(*, (action_dims,)), (*)`, or `None`
- output: `(*)`

value_td_step(trajectory)

The loss function for training the value network using the TD loss. If a_prime is provided, the action-value function Q(s, a) is used. Otherwise, the state-value function V(s) is used.

$$\begin{aligned} \text{LHS} &= V(s) \\ \text{RHS} &= \begin{cases} r + V(s') & \text{if not terminal} \\ r & \text{if terminal} \end{cases} \end{aligned}$$

The LHS and RHS can be optimized using `MSELoss` or `SmoothL1Loss`.

Parameters

- `trajectory` (`Trajectory`) – A batch of trajectories.
- `a_prime` (`torch.Tensor` / `None`) – The next action tensor. If provided, the action-value function $Q(s, a)$ is used. If not provided, the state-value function $V(s)$ is used.

Returns

The LHS and RHS of the TD loss.

Return type

`Tuple[torch.Tensor, torch.Tensor]`

```
class utils.ctorch.rl.CircularTensor(size, dtype=torch.float32)
```

Bases: `Module`

Implements a circular buffer for storing tensors. Supports `state_dict()` and `load_state_dict()` for checkpointing.

Parameters

- `size` (`int`) – The maximum number of elements in the buffer.
- `dtype` (`torch.dtype`, `optional`) – The data type of the elements in the buffer. Defaults to `torch.float32`.

`append(value)`

Append a batch of new elements to the buffer.

Parameters

`value` (`Any`) – The new elements to append. Supports any input that can be converted to a tensor.

Shapes:

- Input shape: $(B, *)$

`as_numpy(force=False)`

Get the underlying numpy array.

`as_tensor()`

Get the underlying tensor.

`forward(x)`

Get a batch of elements according to the index.

`property size: int`

Get the maximum size of the buffer.

Returns

The maximum size of the buffer.

Return type

`int`

```
class utils.ctorch.rl.PrioritizedReplayBuffer(size, continuous_action=False, p_max=1000.0)
```

Bases: `ReplayBuffer`

Implements a prioritized replay buffer for storing and sampling experiences.

Parameters

- **size** (*int*) – The maximum size of the buffer.
- **continuous_action** (*bool, optional*) – Whether the action space is continuous. Defaults to False.
- **p_max** (*float / int, optional*) – Default priority value for newly incoming experiences. Defaults to 1e3.

get_ipw(*idx, beta=1.0, eps=0.0005*)

Get the inverse probability weighting (IPW) for a batch of experiences, to balance the loss function. The IPW weight is given by:

$$\tilde{w}_i = \left(\frac{N \cdot w_i}{\sum_j w_j} \right)^\beta$$

Parameters

- **idx** (*torch.Tensor*) – The indices of the experiences to retrieve.
- **beta** (*float*) – The beta parameter for the importance sampling.

Returns

The importance sampling weights for the specified experiences.

Return type

torch.Tensor

get_weight(*idx*)

Get the sampling weight for a batch of experiences.

Return type

Tensor

sample_index(*batch_size*)

Sample a batch of indices from the replay buffer.

Return type

Tensor

set_weight(*idx, weight, alpha=1.0*)

Set the sampling weight for a batch of experiences.

store(*trajectory*)

Store a batch of new experience in the buffer.

class utils.ctrch.rl.ReplayBuffer(*size, continuous_action=False*)

Bases: *Module*

Implements a replay buffer for storing and sampling experiences.

Parameters

- **size** (*int*) – The maximum size of the buffer.
- **continuous_action** (*bool, optional*) – Whether the action space is continuous. Defaults to False.

forward(*idx*)

Sample a batch of experiences from the buffer.

Return type

Trajectory

get_batch(idx)

Sample a batch of experiences from the buffer.

Return type

Trajectory

property length

Get the current length of the buffer.

sample_index(batch_size)

Randomly sample a batch of indices from the buffer.

Return type

Tensor

store(trajectory)

Store a batch of new experience in the buffer.

class utils.ctorch.rl.TargetNetworkMixin

Bases: *Module*

Mixin class for models with a target network.

copy_params(src, tgt, weight=1.0)**setup_target()**

Create a target network by copying the current network. This method can only be called once.

property target: Self

The target network for the Q-learning algorithm. Used in double DQN. If not set, the current network is used.

Returns

The target network or self if not set.

Return type

Module

update_target(weight=None, tau=None)

Update the target network by copying the weights from the current network. No-op if the target network is not used.

Parameters

weight (*float*, *optional*) – The interpolation weight for the update. By default, it is 1.0.

class utils.ctorch.rl.Trajectory(state, action, reward, next_state, done, log_pi, total_reward)

Bases: *object*

A batch of states, actions, and rewards collected from the environment.

Parameters

- **state** (*torch.Tensor*) – The states in the trajectory of shape (L, *state_shape).
- **action** (*torch.Tensor*) – The actions taken in the trajectory of shape (L, *action_shape).
- **reward** (*torch.Tensor*) – The rewards received in the trajectory of shape (L,).
- **next_state** (*torch.Tensor*) – The next states in the trajectory of shape (L, *state_shape).
- **done** (*torch.Tensor*) – The done flags in the trajectory of shape (L,).

- **log_pi** (`torch.Tensor`) – The log probabilities of the actions taken in the trajectory of shape (L,).
- **total_reward** (`float`) – The total reward of the trajectory. `nan` if the batch is not a trajectory.

action: `Tensor`

clone()

done: `Tensor`

get(key)

Return type

`Tensor`

log_pi: `Tensor`

next_state: `Tensor`

reward: `Tensor`

state: `Tensor`

tensor_fields: `ClassVar[List[str]]` = ['state', 'action', 'reward', 'next_state', 'done', 'log_pi']

total_reward: `float`

`utils.ctorch.rl.run_episode(env, model, max_episode_steps=None, reward_shape=<function _default_shape>, **act_kwargs)`

Simulates a single episode in the environment using the given model.

Parameters

- **env** (`gymnasium.Env`) – The environment to simulate.
- **model** (`BaseRLModel`) – The RL model to use for action selection.
- **eps** (`float`) – The exploration rate.
- **max_episode_steps** (`int / None`) – The maximum number of steps in the episode.
- **reward_shape** (`Callable[[s, a, r, s_prime, terminated, truncated], torch.Tensor]`) – A function to adjust the reward based on the current reward, arrived state, terminal state, and time limit truncation. Returns the adjusted reward to be passed to the model.

Returns

The collected trajectory.

Return type

`Trajectory`

`utils.ctorch.rl.torch_step(env, device='cpu')`

A wrapper for the environment step function to convert actions from torch tensors and results to torch tensors.

Parameters

env (`gymnasium.Env`) – The environment to wrap.

Returns

The wrapped `env.step` function.

Return type

Callable[[torch.Tensor], Tuple[torch.Tensor, torch.Tensor, torch.Tensor, torch.Tensor, Dict[str, Any]]]

3.8 Utilities for Uplift Modeling

uplift.py - Module for uplift modeling functionalities in the ctorch library.

`utils.ctorch.uplift.perfect_qini_curve(y_true, treatment, propensity=None)`

Compute the perfect (optimum) Qini curve.

This is a function, given points on a curve. For computing the area under the Qini Curve, see `qini_auc_score()`.

Parameters

- **y_true** (*1d array-like*) – Correct (true) binary target values.
- **treatment** (*1d array-like*) – Treatment labels.
- **propensity** (*1d array-like, optional*) – Propensity scores. Will be inferred as uniform if not provided.

Returns

Points on the perfect Qini curve.

Return type

Tuple[torch.Tensor, torch.Tensor]

`utils.ctorch.uplift.perfect_uplift(y_true, treatment)`

Compute the perfect (optimum) uplift predictions.

Parameters

- **y_true** (*torch.Tensor*) – Correct (true) binary target values.
- **treatment** (*torch.Tensor*) – Treatment labels.

Returns

Perfect uplift predictions.

Return type

torch.Tensor

`utils.ctorch.uplift.perfect_uplift_curve(y_true, treatment, propensity=None)`

Compute the perfect (optimum) Uplift curve.

This is a function, given points on a curve. For computing the area under the Uplift Curve, see `uplift_auc_score()`.

Parameters

- **y_true** (*1d array-like*) – Correct (true) binary target values.
- **treatment** (*1d array-like*) – Treatment labels.
- **propensity** (*1d array-like, optional*) – Propensity scores. Will be inferred as uniform if not provided.

Returns

Points on the perfect Uplift curve.

Return type

Tuple[torch.Tensor, torch.Tensor]

`utils.ctorch.uplift.qini_auc_score(y_true, uplift, treatment, propensity=None)`

Compute normalized Area Under the Qini Curve from prediction scores.

By computing the area under the Qini curve, the curve information is summarized in one number. For binary outcomes the ratio of the actual uplift gains curve above the diagonal to that of the optimum Qini Curve.

Parameters

- **y_true** (*1d array-like*) – Correct (true) binary target values.
- **uplift** (*1d array-like*) – Predicted uplift, as returned by a model.
- **treatment** (*1d array-like*) – Treatment labels.
- **propensity** (*1d array-like, optional*) – Propensity scores. Will be inferred as uniform if not provided.

Returns

Normalized Area Under the Qini Curve.

Return type

float

`utils.ctorch.uplift.qini_curve(y_true, uplift, treatment, propensity=None)`

Compute Qini curve.

For computing the area under the Qini Curve, see `qini_auc_score()`.

Parameters

- **y_true** (*1d array-like*) – Correct (true) binary target values.
- **uplift** (*1d array-like*) – Predicted uplift, as returned by a model.
- **treatment** (*1d array-like*) – Treatment labels.
- **propensity** (*1d array-like, optional*) – Propensity scores. If provided, inverse propensity score weighting is applied.

Returns

x and y points on the Qini curve.

Return type

Tuple[torch.Tensor, torch.Tensor]

`utils.ctorch.uplift.uplift_auc_score(y_true, uplift, treatment, propensity=None)`

Compute normalized Area Under the Uplift Curve from prediction scores.

By computing the area under the Uplift curve, the curve information is summarized in one number. For binary outcomes the ratio of the actual uplift gains curve above the diagonal to that of the optimum Uplift Curve.

Parameters

- **y_true** (*1d array-like*) – Correct (true) binary target values.
- **uplift** (*1d array-like*) – Predicted uplift, as returned by a model.
- **treatment** (*1d array-like*) – Treatment labels.
- **propensity** (*1d array-like, optional*) – Propensity scores. Will be inferred as uniform if not provided.

Returns

Normalized Area Under the Uplift Curve.

Return type

float

`utils.ctorch.uplift.uplift_curve(y_true, uplift, treatment, propensity=None)`

Compute Uplift curve.

For computing the area under the Uplift Curve, see `uplift_auc_score()`.**Parameters**

- **y_true** (*1d array-like*) – Correct (true) binary target values.
- **uplift** (*1d array-like*) – Predicted uplift, as returned by a model.
- **treatment** (*1d array-like*) – Treatment labels.
- **propensity** (*1d array-like, optional*) – Propensity scores. Will be inferred as uniform if not provided.

Returns

x and y points on the Uplift curve.

Return type

Tuple[torch.Tensor, torch.Tensor]

`utils.ctorch.uplift.validate_binary(*args)`

Validate that all input tensors are binary (contain only 0s and 1s).

Parameters***args** (*torch.Tensor*) – Tensors to validate.**Raises****ValueError** – If any tensor contains values other than 0 or 1.**Return type**

None

`utils.ctorch.uplift.validate_device(*args)`

Validate that all input tensors are on the same device.

Parameters***args** (*torch.Tensor*) – Tensors to validate.**Returns**

The common device of the input tensors.

Return type

torch.device

Raises**ValueError** – If the tensors are on different devices.`utils.ctorch.uplift.validate_shape(*args, batch_dimension=0)`

Validate that all input tensors have the same shape on the batch dimension.

Parameters

- ***args** (*torch.Tensor*) – Tensors to validate.
- **batch_dimension** (*int*) – The dimension representing the batch size.

Returns

The size of the batch dimension.

Return type

int

Raises

ValueError – If the shapes of the tensors do not match in the batch dimension.

`utils.ctorch.uplift.validate_tensor(*args, dtype=torch.float32)`

Convert input arguments to PyTorch tensors of a specified data type. :type *args: Tensor :param *args: Input arguments to convert. :type dtype: dtype :param dtype: Desired data type for the tensors. :type dtype: torch.dtype

Raises

TypeError – If any argument cannot be converted to a tensor.

Return type

List[Tensor]

**CHAPTER
FOUR**

SAMPLER UTILITIES

LOGGING UTILITIES

`utils.clogging` - Colored Logging Formatter

A Python module that provides a colored logging formatter for the `logging` module. This module defines a base colored formatter class and a default implementation that applies specific colors to different logging levels and fields. It allows for easy customization of log message styles, making it easier to distinguish between different log levels in console output.

```
class utils.clogging.BaseColoredFormatter(fmt='%(asctime)s %(levelname)-8s %(name)s: %(message)s',
                                            datefmt='%Y-%m-%d %H:%M:%S', style='%')
```

Bases: `Formatter`

Add color to log messages based on their level.

This formatter allows customization of log message styles based on their logging level. It supports coloring the level name and optionally the message itself. The styles can be customized through the `get_color` method, which should return a argument dictionary compatible with `cprint.cprint`.

Parameters

- `fmt (str)` – The format string for the log messages.
- `datefmt (str / None)` – The format string for the date in log messages.
- `style (str)` – The character used in the format string (default is ‘%’).

(same as `logging.Formatter`)

`property field_names: List[str]`

Returns the list of field names that can be colored. This can be overridden in subclasses to specify which fields are available.

‘base’ is a special field that can be used to apply a default color to the template.

Returns

A list of field names that can be colored.

Return type

`List[str]`

`format(record)`

Formats the log record with colors applied to the specified fields.

Parameters

`record (logging.LogRecord)` – The log record to format.

Returns

The formatted log message with colors applied.

Return type

str

formatTime(record, datefmt=None)

Formats the time of the log record. If datefmt is provided, it formats the time accordingly. Otherwise, it uses the default format.

Parameters

- **record** (*logging.LogRecord*) – The log record to format.
- **datefmt** (*str* / *None*) – The format string for the date.

Returns

The formatted time string with color applied.

Return type

str

get_color(field_name, level)

Returns the color format for a given field name and logging level. This method should be overridden in subclasses to provide specific color styles for different fields and levels.

Parameters

- **field_name** (*str*) – The name of the field to color.
- **level** (*int*) – The logging level for which to get the color.

Returns

A dictionary of color attributes compatible with *utils.cprint.cprint*.

Return type

Dict[str, Any] | None

```
class utils.clogging.DefaultColoredFormatter(fmt='%(asctime)s %(levelname)-8s %(name)s: %(message)s', datefmt='%Y-%m-%d %H:%M:%S', style='%')
```

Bases: *BaseColoredFormatter*

Default colored formatter with predefined styles for log levels.

get_color(field_name, level)

Returns the color format for a given field name and logging level. This method should be overridden in subclasses to provide specific color styles for different fields and levels.

Parameters

- **field_name** (*str*) – The name of the field to color.
- **level** (*int*) – The logging level for which to get the color.

Returns

A dictionary of color attributes compatible with *utils.cprint.cprint*.

Return type

Dict[str, Any] | None

PRINTING UTILITIES

utils.cprint - Colorful Terminal Output

A Python module for colorful terminal output with support for ANSI escape codes, xterm-256color, and true color (24-bit RGB). It provides functions to format text with foreground and background colors, styles (bold, italic, underline, strikethrough), and fallback to lower color grades if the terminal does not support the requested color depth.

```
utils.cprint.cformat(text, *, fg=None, bg=None, fallback=True, bf=False, dim=False, it=False, us=False, st=False, reset='\x1b[0m')
```

Format a string with ANSI escape codes for colors and styles.

Parameters

- **text (str)** – The text to format.
- **fg (str / int / Tuple[int, int, int] / None)** – Foreground color (color name, xterm-256color index, RGB tuple, or hex code).
- **bg (str / int / Tuple[int, int, int] / None)** – Background color (same format as fg).
- **fallback (bool)** – Whether to use fallback colors if the terminal does not support the requested color depth.
- **bf (bool)** – Whether to use bold text.
- **it (bool)** – Whether to use italic text.
- **us (bool)** – Whether to underline the text.
- **st (bool)** – Whether to use strikethrough text (not supported in all terminals).
- **reset (str)** – String to append after the formatted text (default is ANSI reset code).

Returns

A formatted string with ANSI escape codes.

Return type

str

```
utils.cprint.cprefix(fg=None, bg=None, fallback=True, bf=False, dim=False, it=False, us=False, st=False)
```

Generate an ANSI escape code prefix for formatting text with colors and styles.

Parameters

- **fg (str / int / Tuple[int, int, int] / None)** – Foreground color (color name, xterm-256color index, RGB tuple, or hex code).
- **bg (str / int / Tuple[int, int, int] / None)** – Background color (same format as fg).

- **fallback** (*bool*) – Whether to use fallback colors if the terminal does not support the requested color depth.
- **bf** (*bool*) – Whether to use bold text.
- **it** (*bool*) – Whether to use italic text.
- **us** (*bool*) – Whether to underline the text.
- **st** (*bool*) – Whether to use strikethrough text (not supported in all terminals).

Returns

A string containing the ANSI escape code prefix.

Return type

str

```
utils.cprint.cprint(*obj, fg=None, bg=None, fallback=True, bf=False, dim=False, it=False, us=False, st=False, reset='\x1b[0m', sep=' ', end='\n', file=None, flush=False)
```

Colorful print function with support for foreground and background colors, styles (bold, italic, underline, strikethrough), and fallback to lower color grades if the terminal does not support the requested color depth.

Parameters

- ***obj** (*Any*) – Objects to print.
- **fg** (*str* / *int* / *Tuple[int, int, int]* / *None*) – Foreground color (color name, xterm-256color index, RGB tuple, or hex code).
- **bg** (*str* / *int* / *Tuple[int, int, int]* / *None*) – Background color (same format as fg).
- **fallback** (*bool*) – Whether to use fallback colors if the terminal does not support the requested color depth.
- **bf** (*bool*) – Whether to use bold text.
- **it** (*bool*) – Whether to use italic text.
- **us** (*bool*) – Whether to underline the text.
- **st** (*bool*) – Whether to use strikethrough text (not supported in all terminals).
- **reset** (*str*) – String to append after the formatted text (default is ANSI reset code).

The following parameters are inherited from the built-in print function:

- **sep**: Separator between objects.
- **end**: String appended after the last object.
- **file**: A file-like object (default is sys.stdout).
- **flush**: Whether to forcibly flush the output buffer.

ARGUMENT PARSING UTILITIES

`utils.parser`

A utility to automatically parse command line arguments into a dataclass instance. This module provides a decorator `@auto_cli` that can be applied to a dataclass. After applying the decorator, you can call `parse_args()` on the dataclass to parse command line arguments and return an instance of the dataclass.

1. A special field `--config` is added to allow loading configuration from a file in JSON, YAML, or TOML format.
2. When a field is specified in both the command line arguments and the configuration file, the command line argument takes precedence.

`utils.parser.auto_cli(cls=None, / (Positional-only parameter separator (PEP 570)), **decorator_kw)`

Automatically generates an argument parser for a dataclass with type hints. Field types, including basic types, and `Optional` are automatically recognized and converted corresponding argument types. Complex types like lists and dictionaries can be input as strings in json format.

Apart from the fields of the dataclass, `auto_cli` also adds a `--config` argument to the parser for loading configuration settings from a json, yaml, or toml file. Values from the configuration file have a lower priority than command line arguments.

The decorated function will get the following new class methods:

- `get_parser(prefix: str = '') -> argparse.ArgumentParser:`

Returns an `argparse.ArgumentParser` instance with arguments based on the dataclass fields. The `prefix` argument is prepended to each argument name. The parser returned does not include the `--config` argument for loading configuration files.

Example:

```
@auto_cli
@dataclasses.dataclass
class YourClass:
    name: str = 'DefaultName'
    age: int = 25

parser = YourClass.get_parser()
parser_prefix = YourClass.get_parser(prefix='man')
```

The generated `parser` will accept the following arguments:

- `--name`: The name argument with a default value of ‘DefaultName’.
- `--age`: The age argument with a default value of 25.

The generated `parser_prefix` will accept the following arguments:

- `--man-name`: The name argument with a default value of ‘DefaultName’.

- `--man-age`: The age argument with a default value of 25.
- `parse_namespace(ns: argparse.Namespace, kw: Dict[str, Any] = None, prefix: str = '') -> 'YourClass'`:
Parses an `argparse.Namespace` instance and a `kwargs` dictionary into an instance of the dataclass. Prefix is prepended to each argument name. The `kw` dictionary is loaded from a config file if specified in `ns`. A `ValueError` is raised if a required argument is missing from both `ns` and `kw`, and no default value is provided.
- `parse_args(argv: List[str] = None) -> 'YourClass'`:
Parses command line arguments and returns an instance of the dataclass. If `argv` is `None`, it uses `sys.argv[1:]`. The method also handles config files specified in the command line arguments.

Returns

The decorated class.

Return type

`Type[cls]`

`utils.parser.get_all_parser(dataclass=None, **dataclasses)`

Returns a combined `argparse.ArgumentParser` that merges the parsers of multiple dataclasses into one, and then adds a `--config` argument for loading configuration files.

The dataclass provided as a positional argument is treated as unprefixed, while the others are prefixed with their keyword argument names.

Parameters

- `dataclass (Optional[Any])` – A single dataclass to include in the parser.
- `**dataclasses (Any)` – Additional dataclasses to include in the parser.

Returns

An argument parser that includes all specified dataclasses.

Return type

`argparse.ArgumentParser`

`utils.parser.parse_all_args(cli_args=None, dataclass=None, **dataclasses)`

Parse command line arguments into a dictionary of dataclass instances. Each dataclass is identified by its name in the `dataclasses` argument.

The `parse_all_args` function accepts two types of input:

1. `parse_all_args(cli_args, dataclass, **dataclasses)`: where `cli_args` is a list of command line arguments, `dataclass` is the unprefixed dataclass, and `dataclasses` are additional prefixed dataclasses.
2. `parse_all_args(dataclass, **dataclasses)`: where `dataclass` is the unprefixed dataclass and `dataclasses` are additional prefixed dataclasses. `sys.argv[1:]` is used as the command line arguments.

The result is a dictionary where the keys are the prefixes of the dataclasses and the values are instances of those dataclasses, parsed from the command line arguments. The unprefixed dataclass is stored under the key ''.

Example:

```
@auto_cli
@dataclasses.dataclass
class ClassMain:
    name: str = 'MainName'
    age: int = 30
```

(continues on next page)

(continued from previous page)

```
@auto_cli
@dataclasses.dataclass
class ClassAdditional:
    address: str = 'DefaultAddress'
    phone: str = '1234567890'

parser = get_all_parser(ClassMain, additional=ClassAdditional)
```

The generated parser will accept the following arguments:

- `--name`: The name argument with a default value of '`MainName`'.
- `--age`: The age argument with a default value of `30`.
- `--additional-address`: The address argument with a default value of '`DefaultAddress`'.
- `--additional-phone`: The phone argument with a default value of '`1234567890`'.

Or the following configuration files:

```
name: 'MainName'
age: 30
additional_address: 'DefaultAddress'
additional_phone: '1234567890'
```

Parameters

- `cli_args (List[str] / Any / None)` – Command line arguments to parse. If `None`, uses `sys.argv[1:]`.
- `dataclass (Any / None)` – A single dataclass to include in the parsing.
- `**dataclasses (Any)` – Additional dataclasses to include in the parsing.

Returns

A dictionary where keys are dataclass names and values are instances of those dataclasses.

Return type

`Dict[str, Any]`

PYTHON MODULE INDEX

U

utils, ??
utils.clogging, 49
utils.cprint, 51
utils.ctorch, 5
utils.ctorch.device, 5
utils.ctorch.functional, 6
utils.ctorch.metrics, 9
utils.ctorch.nn, 11
utils.ctorch.ops, 24
utils.ctorch.padding, 25
utils.ctorch.rl, 30
utils.ctorch.uplift, 43
utils.parser, 53
utils.sampler, 47

INDEX

A

`A()` (*utils.ctorch.rl.BaseValueNetwork method*), 36
`A_all()` (*utils.ctorch.rl.BaseQNetwork method*), 34
`accumulator()` (*utils.ctorch.metrics.BatchedAUC method*), 9
`accumulator()` (*utils.ctorch.metrics.BatchedHitRate method*), 9
`accumulator()` (*utils.ctorch.metrics.BatchedMetric method*), 9
`accumulator()` (*utils.ctorch.metrics.BatchedNDCG method*), 10
`act()` (*utils.ctorch.rl.BasePolicyNetwork method*), 32
`act()` (*utils.ctorch.rl.BaseQNetwork method*), 35
`act()` (*utils.ctorch.rl.BaseRLModel method*), 36
`act()` (*utils.ctorch.rl.BaseValueNetwork method*), 37
`action` (*utils.ctorch.rl.Trajectory attribute*), 42
`action_td_step()` (*utils.ctorch.rl.BaseValueNetwork method*), 38
`Activation` (*class in utils.ctorch.nn*), 11
`append()` (*utils.ctorch.rl.CircularTensor method*), 39
`append_right()` (*in module utils.ctorch.padding*), 25
`as_numpy()` (*utils.ctorch.rl.CircularTensor method*), 39
`as_tensor()` (*utils.ctorch.rl.CircularTensor method*), 39
`auc_score()` (*in module utils.ctorch.metrics*), 10
`auto_cli()` (*in module utils.parser*), 53
`auuc_score()` (*in module utils.ctorch.metrics*), 10
`avg_free_gb` (*utils.ctorch.device.GpuStat attribute*), 5
`avg_util` (*utils.ctorch.device.GpuStat attribute*), 5

B

`backward()` (*utils.ctorch.ops.GradientReversalOp static method*), 24
`BaseColoredFormatter` (*class in utils.clogging*), 49
`BaseDistributionalQNetwork` (*class in utils.ctorch.rl*), 30
`BasePolicyNetwork` (*class in utils.ctorch.rl*), 32
`BaseQNetwork` (*class in utils.ctorch.rl*), 33
`BaseRLModel` (*class in utils.ctorch.rl*), 36
`BaseValueNetwork` (*class in utils.ctorch.rl*), 36
`BatchedAUC` (*class in utils.ctorch.metrics*), 9
`BatchedHitRate` (*class in utils.ctorch.metrics*), 9
`BatchedMetric` (*class in utils.ctorch.metrics*), 9

`BatchedNDCG` (*class in utils.ctorch.metrics*), 9

C

`cformat()` (*in module utils.cprint*), 51
`check_input_shape()` (*utils.ctorch.nn.ZeroInflatedLogNormal method*), 24
`CholeskyTrilLinear` (*class in utils.ctorch.nn*), 11
`CircularTensor` (*class in utils.ctorch.rl*), 39
`clone()` (*utils.ctorch.rl.Trajectory method*), 42
`copy_params()` (*utils.ctorch.rl.TargetNetworkMixin method*), 41
`cov()` (*utils.ctorch.nn.GaussianLinear method*), 16
`cprefix()` (*in module utils.cprint*), 51
`cprint()` (*in module utils.cprint*), 52
`cumulative_reward()` (*utils.ctorch.rl.BasePolicyNetwork method*), 32

D

`debug()` (*utils.ctorch.nn.Module method*), 18
`DeEmbedding` (*class in utils.ctorch.nn*), 13
`DefaultColoredFormatter` (*class in utils.clogging*), 50
`device` (*utils.ctorch.nn.Module property*), 18
`diag()` (*utils.ctorch.nn.CholeskyTrilLinear method*), 12
`DNN` (*class in utils.ctorch.nn*), 13
`done` (*utils.ctorch.rl.Trajectory attribute*), 42

E

`exploit()` (*utils.ctorch.rl.BaseQNetwork method*), 35
`explore()` (*utils.ctorch.rl.BaseQNetwork method*), 35

F

`f()` (*utils.ctorch.nn.FactorizedNoisyLinear static method*), 14
`FactorizedNoisyLinear` (*class in utils.ctorch.nn*), 14
`FeatureEmbedding` (*class in utils.ctorch.nn*), 14
`field_names` (*utils.clogging.BaseColoredFormatter property*), 49
`finalize()` (*utils.ctorch.metrics.BatchedMetric method*), 9
`format()` (*utils.clogging.BaseColoredFormatter method*), 49

G

- `GaussianLinear` (*class in utils.ctorch.nn*), 15
- `get()` (*utils.ctorch.rl.Trajectory method*), 42
- `get_all_parser()` (*in module utils.parser*), 54
- `get_attention_map()`
 (*utils.ctorch.nn.TransformerEncoderLayer method*), 23
- `get_batch()` (*utils.ctorch.rl.ReplayBuffer method*), 40
- `get_best_device()` (*in module utils.ctorch.device*), 5
- `get_color()` (*utils.clogging.BaseColoredFormatter method*), 50

H

- `hit_rate()` (*in module utils.ctorch.metrics*), 10

I

- `idx` (*utils.ctorch.device.GpuStat attribute*), 5
- `IndependentNoisyLinear` (*class in utils.ctorch.nn*), 17

L

- `length` (*utils.ctorch.rl.ReplayBuffer property*), 41
- `log_norm_pdf()` (*in module utils.ctorch.functional*), 6
- `log_pi` (*utils.ctorch.rl.Trajectory attribute*), 42
- `log_pi()` (*utils.ctorch.rl.BasePolicyNetwork method*), 33
- `logit_product()` (*in module utils.ctorch.functional*), 6
- `LogStackedTruncatedNormal` (*class in utils.ctorch.nn*), 17
- `loss()` (*utils.ctorch.nn.LogStackedTruncatedNormal method*), 18
- `loss()` (*utils.ctorch.nn.NegativeBinomial method*), 20
- `loss()` (*utils.ctorch.nn.StackedTruncatedNormal method*), 22
- `loss()` (*utils.ctorch.nn.ZeroInflatedLogNormal method*), 24
- `loss()` (*utils.ctorch.rl.BaseRLModel method*), 36

M

- `masked_select()` (*in module utils.ctorch.padding*), 26
- `mean()` (*utils.ctorch.nn.GaussianLinear method*), 16

`MetricFormatter` (*class in utils.ctorch.metrics*), 10
`mmd_distance()` (*in module utils.ctorch.functional*), 6
`module`
 `utils`, 1
 `utils.clogging`, 49
 `utils.cprint`, 51
 `utils.ctorch`, 5
 `utils.ctorch.device`, 5
 `utils.ctorch.functional`, 6
 `utils.ctorch.metrics`, 9
 `utils.ctorch.nn`, 11
 `utils.ctorch.ops`, 24
 `utils.ctorch.padding`, 25
 `utils.ctorch.rl`, 30
 `utils.ctorch.uplift`, 43
 `utils.parser`, 53
 `utils.sampler`, 47
`Module` (*class in utils.ctorch.nn*), 18
`MonotonicLinear` (*class in utils.ctorch.nn*), 18
`MultivariateNormalClass` (*class in utils.ctorch.nn*), 19

N

`name` (*utils.ctorch.device.GpuStat attribute*), 5
`ndcg_score()` (*in module utils.ctorch.metrics*), 11
`NegativeBinomial` (*class in utils.ctorch.nn*), 19
`next_state` (*utils.ctorch.rl.Trajectory attribute*), 42
`norm_pdf()` (*in module utils.ctorch.functional*), 7
`normalize_trajectory()`
 (*utils.ctorch.rl.BasePolicyNetwork method*), 33
`num_parameters` (*utils.ctorch.nn.Module property*), 18

P

`pack_padded_sequence_right()` (*in module utils.ctorch.padding*), 27
`packed_binary_op()` (*in module utils.ctorch.padding*), 27
`packed_concat()` (*in module utils.ctorch.padding*), 27
`packed_forward()` (*in module utils.ctorch.padding*), 28
`packed_unary_op()` (*in module utils.ctorch.padding*), 28
`pad_packed_sequence_right()` (*in module utils.ctorch.padding*), 28
`parse_all_args()` (*in module utils.parser*), 54
`pd()` (*utils.ctorch.nn.CholeskyTrilLinear method*), 12
`perfect_qini_curve()` (*in module utils.ctorch.uplift*), 43
`perfect_uplift()` (*in module utils.ctorch.uplift*), 43
`perfect_uplift_curve()` (*in module utils.ctorch.uplift*), 43
`pi()` (*utils.ctorch.rl.BasePolicyNetwork method*), 33
`policy_parameters()`
 (*utils.ctorch.rl.BasePolicyNetwork method*), 33
`prepend_left()` (*in module utils.ctorch.padding*), 29

`PrioritizedReplayBuffer` (*class in utils.ctorch.rl*), 39

Q

`Q()` (*utils.ctorch.rl.BaseDistributionalQNetwork method*), 30
`Q()` (*utils.ctorch.rl.BaseValueNetwork method*), 37
`Q_all()` (*utils.ctorch.rl.BaseDistributionalQNetwork method*), 31
`Q_all()` (*utils.ctorch.rl.BaseQNetwork method*), 34
`Q_all_dist()` (*utils.ctorch.rl.BaseDistributionalQNetwork method*), 31
`Q_dist()` (*utils.ctorch.rl.BaseDistributionalQNetwork method*), 31
`qini_auc_score()` (*in module utils.ctorch.uplift*), 44
`qini_curve()` (*in module utils.ctorch.uplift*), 44

R

`rbf_kernel()` (*in module utils.ctorch.functional*), 7
`ReplayBuffer` (*class in utils.ctorch.rl*), 40
`reset()` (*utils.ctorch.metrics.BatchedMetric method*), 9
`reward` (*utils.ctorch.rl.Trajectory attribute*), 42
`RotaryTemporalEmbedding` (*class in utils.ctorch.nn*), 20
`run_episode()` (*in module utils.ctorch.rl*), 42

S

`sample_index()` (*utils.ctorch.rl.PrioritizedReplayBuffer method*), 40
`sample_index()` (*utils.ctorch.rl.ReplayBuffer method*), 41
`set_weight()` (*utils.ctorch.rl.PrioritizedReplayBuffer method*), 40
`setup_target()` (*utils.ctorch.rl.BasePolicyNetwork method*), 33
`setup_target()` (*utils.ctorch.rl.TargetNetworkMixin method*), 41
`SinusoidalTemporalEmbedding` (*class in utils.ctorch.nn*), 21
`size` (*utils.ctorch.rl.CircularTensor property*), 39
`StackedTruncatedNormal` (*class in utils.ctorch.nn*), 21
`state` (*utils.ctorch.rl.Trajectory attribute*), 42
`store()` (*utils.ctorch.rl.PrioritizedReplayBuffer method*), 40
`store()` (*utils.ctorch.rl.ReplayBuffer method*), 41

T

`target` (*utils.ctorch.rl.TargetNetworkMixin property*), 41
`TargetNetworkMixin` (*class in utils.ctorch.rl*), 41
`td_step()` (*utils.ctorch.rl.BaseQNetwork method*), 36
`tensor_fields` (*utils.ctorch.rl.Trajectory attribute*), 42
`torch_step()` (*in module utils.ctorch.rl*), 42
`total_embedding_size`
 (*utils.ctorch.nn.FeatureEmbedding property*), 15

`total_gb (utils.ctorch.device.GpuStat attribute), 5`
`total_reward (utils.ctorch.rl.Trajectory attribute), 42`
`Trajectory (class in utils.ctorch.rl), 41`
`TransformerDecoderLayer (class in utils.ctorch.nn), 22`
`TransformerEncoderLayer (class in utils.ctorch.nn), 23`
`truncate_left() (in module utils.ctorch.padding), 29`
`truncate_right() (in module utils.ctorch.padding), 29`

U

`unpad_sequence_right() (in module utils.ctorch.padding), 30`
`update() (utils.ctorch.metrics.MetricFormatter method), 10`
`update_target() (utils.ctorch.rl.TargetNetworkMixin method), 41`
`uplift_auc_score() (in module utils.ctorch.uplift), 44`
`uplift_curve() (in module utils.ctorch.uplift), 45`
`utils`
 `module, 1`
`utils.clogging`
 `module, 49`
`utils.cprint`
 `module, 51`
`utils.ctorch`
 `module, 5`
`utils.ctorch.device`
 `module, 5`
`utils.ctorch.functional`
 `module, 6`
`utils.ctorch.metrics`
 `module, 9`
`utils.ctorch.nn`
 `module, 11`
`utils.ctorch.ops`
 `module, 24`
`utils.ctorch.padding`
 `module, 25`
`utils.ctorch.rl`
 `module, 30`
`utils.ctorch.uplift`
 `module, 43`
`utils.parser`
 `module, 53`
`utils.sampler`
 `module, 47`

V

`V() (utils.ctorch.rl.BaseQNetwork method), 34`
`V() (utils.ctorch.rl.BaseValueNetwork method), 37`
`V_dist() (utils.ctorch.rl.BaseDistributionalQNetwork method), 32`
`validate_binary() (in module utils.ctorch.uplift), 45`

`validate_device() (in module utils.ctorch.uplift), 45`
`validate_shape() (in module utils.ctorch.uplift), 45`
`validate_tensor() (in module utils.ctorch.uplift), 46`
`value_model (utils.ctorch.rl.BasePolicyNetwork property), 33`
`value_parameters() (utils.ctorch.rl.BasePolicyNetwork method), 33`
`value_td_step() (utils.ctorch.rl.BaseValueNetwork method), 38`

W

`wasserstein_distance() (in module utils.ctorch.functional), 8`

Z

`ZeroInflatedLogNormal (class in utils.ctorch.nn), 23`